

Introduction to Configuration Management

for Linux Clusters Institute Introductory Workshop

Mike Renfro, PhD

Information Technology Services, Tennessee Tech University

2025-02-11

Acknowledgments

- ▶ Thanks to Alexei Kotelnikov from Rutgers University

Stateful

- ▶ OS, settings, some/all applications for compute nodes stored on local disks
- ▶ Local files are preserved across reboots
- ▶ Primary challenge: **keeping things consistent across nodes**
- ▶ Common tools: **Ansible, Chef, Puppet, Salt**
- ▶ Model used in LCI workshop labs

Stateless

- ▶ OS, applications, settings for compute nodes pulled from central network server
- ▶ Root filesystem is a ramdisk, read-only NFS mount, or local disk
- ▶ Local files are **not** preserved across reboots
- ▶ Primary challenges: network traffic, load on NFS server, managing the contents of the root filesystem contents
- ▶ Common tools: **initramfs-tools**, **dracut**, **xCAT**, **Warewulf**
- ▶ Model used by OpenHPC project (and others)

Defining configuration management

The management and maintenance of operating systems, applications, and services via automation instead of manual intervention.

More formally:

- ▶ **declaring system state in a repeatable and auditable fashion**
- ▶ **using tools to impose state and keep systems from deviating**

Adjacent to, but not identical to, “orchestration” (Terraform, Kubernetes, etc.)

Benefits of configuration management

- ▶ Centralized, authoritative catalog of system configuration
- ▶ Automated enforcement of system state
- ▶ Ensured consistency among systems
- ▶ Mix and match components
- ▶ Collection of system “facts” to inform decision making
- ▶ Usually tracked with version control (Git or equivalent)

Cautions about configuration management

- ▶ You need some understanding of the more manual processes for systems administration tasks
- ▶ Analogous to needing to know how to calculate a formula on paper before writing a program to calculate the same formula
- ▶ Automation should be used to make things more consistent, not to just blindly copy/paste

Historical: scripts, make (and scripts)

The DevOps folks didn't invent “infrastructure as code” or “defined state”:

- ▶ Systems administrators have been writing shell scripts for 50+ years
- ▶ <http://www.infrastructures.org> has references going back to 1988
- ▶ Clever use of network booting, make, shell commands, and output files can take you a long way

Modern tools

Most modern tools use:

- ▶ a syntax describing desired client system state (usually in a more abstract sense than shell commands)
- ▶ facts the client knows about itself (hostname, OS version, “am I a physical or virtual machine?”, ...)
- ▶ ordering of steps required to reach the desired system state
- ▶ comparison of current system state to desired system state to avoid excess/redundant work (idempotence)

Sampling of prominent tools

- Puppet (2005–)** Ruby-based, with agent installed on clients. Communication via HTTPS to central server (port 8140). Ruby-like domain specific language (DSL) for describing state.
- Chef Infra (2009–)** Ruby-based, with agent installed on clients. Communication via HTTP and HTTPS to central server (ports 443 and 9683). Uses literal Ruby instead of a DSL.
- Salt (2011–)** Python-based, with agent installed on clients. Communication via AES-encrypted ZeroMQ to central server. Uses YAML for its DSL.
- Ansible (2012–)** Python-based, no agent installed on clients. Settings pushed from central server or admin workstation over ssh. Uses YAML for its DSL.

How Ansible works

- ▶ Software installed on any administrative endpoint (`pip install ansible`)
- ▶ Uses public-key or host-based `ssh` authentication to managed clients
- ▶ Runs several `ssh` connections at once for managing multiple clients
- ▶ Elevates privileges with `sudo`
- ▶ Builds custom Python scripts to copy to, and execute on, clients
- ▶ Lower-level tasks can be composed into playbooks, and playbooks can be composed into roles.
- ▶ In HPC, Ansible could be used to configure stateful nodes, or to configure the central server managing stateless nodes.

An Ansible story

- ▶ Half-day tutorial at PEARC25 — “OpenHPC: Beyond the Install Guide”
- ▶ Similar infrastructure requirements to LCI
- ▶ 31 small clusters hosted at NSF’s Jetstream2 OpenStack facility
- ▶ Created nodes, networks, IPs with Terraform
- ▶ Needed to configure each management node (head node) as if we’d just finished installing OpenHPC using instructions from the install guide
- ▶ OpenHPC provides an automated installation script with commands generated from a common source with documentation
- ▶ I’d already run through the manual steps several times
- ▶ I’d already run the provided script several times

Example file structure

For configuring an OpenHPC system management server (head node):

```
% find ansible -type f
ansible/ansible.cfg
ansible/hosts.ini
ansible/0-undocumented-prereqs-unrelated-settings.yaml
ansible/2-install-base-os.yaml
ansible/3-install-openhpc-components.yaml
ansible/a0-installation-template.yaml
ansible/a1-run-recipe.yaml
...
```

ansible.cfg

```
[defaults]
inventory = hosts.ini
remote_user = rocky

[privilege_escalation]
become = true
```

hosts.ini

If an LCI cluster was running OpenHPC:

```
[head]  
lci-head-[01:55]-1.ncsa.cloud
```

Tasks for Excerpt of OpenHPC Install Guide 3.2, Section 2

```
# 2-install-base-os.yaml
- name: 2. Install Base Operating System
  hosts: head
  tasks:
    # [sms]# echo ${sms_ip} ${sms_name} >> /etc/hosts
    - name: Add hostname to /etc/hosts
      ansible.builtin.lineinfile:
        path: /etc/hosts
        line: "172.16.0.1 {{ ansible_hostname }}"
    # [sms]# setenforce 0
    - name: Disable SELinux
      ansible.posix.selinux:
        state: disabled
```


Tasks for Excerpt of OpenHPC Install Guide 3.2, Section 3

(Part 1 of 2)

```
# 3-install-openhpc-components.yaml
- name: 3. Install OpenHPC Components
  hosts: head
  tasks:
    - name: Check if crb repository is enabled
      ansible.builtin.shell: |
        dnf repolist crb | grep -q enabled
      register: crb_enabled
      changed_when: false
      ignore_errors: true
```

Tasks for Excerpt of OpenHPC Install Guide 3.2, Section 3

(Part 2 of 2)

```
- name: Ensure the crb repository is enabled
  ansible.builtin.shell: |
    dnf install -y dnf-plugins-core
    dnf config-manager --set-enabled crb
  when: crb_enabled is failed
```

Running a Play

```
% ansible-playbook 2-install-base-os.yaml
```

Conclusions

- ▶ Configuration management can provide:
 - ▶ better consistency across systems with common attributes
 - ▶ fewer manual errors
 - ▶ less reinventing wheels
- ▶ The specific technology doesn't matter as much as the community support, both locally and globally